Data Science Optimization Frameworks for Computationally Efficient Scientific Simulation: A Cross-Domain Narrative Review

Sabir Z. Sharipov
Independent Researcher
Bukhara, Uzbekistan
snikeee11@gmail.com
ORCID: 0009-0007-4619-1665

March 2026

Working Paper / Preprint — Not yet peer-reviewed

Abstract

Background. Artificial intelligence and scientific simulation workloads consume an estimated 1–2% of global electricity, with continued rapid growth projected. This computational cost concentrates advanced simulation capabilities among well-resourced institutions, creating barriers for academic researchers, startups, and organizations in resource-constrained regions. Data science methodologies have independently demonstrated substantial computational savings across scientific domains, yet no unified framework has synthesized these techniques for cross-domain application.

Objective. This review proposes a four-tier data science optimization framework for energy-efficient scientific simulation, assessing its applicability across five domains: healthcare digital twins, computational chemistry, computational neuroscience, manufacturing optimization, and digital olfaction.

Methods. A structured literature search was conducted across Google Scholar, IEEE Xplore, and arXiv, supplemented by reference list scanning, covering January 2018 through December 2025. Studies were included if they empirically quantified at least one computational efficiency metric alongside accuracy reporting.

Results. Original experiments identify three empirical boundary conditions for the framework's applicability, absent from the existing literature. First, transfer learning produces consistently negative results when source and target domains have different generative processes — observed in two of six proxy datasets and confirmed under controlled physics-based computation (ethanol→methanol transfer using ANI-2x, mean $\Delta R^2 = -0.445$). This distinguishes distribution shift *within* a generative process from shift *between* processes — a distinction underrepresented in the transfer learning literature, which predominantly reports positive results. Second, the framework fails entirely when baseline feature predictability falls below $R^2 > 0.3$ (healthcare dataset, $R^2 = 0.153$), establishing a minimum threshold for applicability. Third, a controlled wall-clock benchmark against the ANI-2x physics potential demonstrates 10^2–$10^3\times$ surrogate speedup with $R^2 = 0.984$, validating speedup claims with measured computation times. Across domains where boundary conditions are met, the four-tier framework — (1) domain-informed feature engineering, (2) surrogate modeling, (3) transfer learning, and (4) model compression — yields computational cost reductions of up to 70–90% under favorable conditions while maintaining 85–95% predictive accuracy. All experiments were conducted using freely available computing resources.

Conclusions. Systematic application of established data science techniques constitutes a generalizable strategy for reducing the computational cost of scientific simulation, with implications for democratizing access to advanced simulation capabilities for resource-constrained institutions worldwide.

Keywords: surrogate modeling; energy-efficient AI; scientific simulation; transfer learning; model compression; computational sustainability; digital twins; physics-informed machine learning

1 Introduction

1.1 The Computational Cost of Scientific Simulation

Scientific simulation underpins modern discovery across medicine, chemistry, materials science, neuroscience, and engineering. However, the computational requirements of high-fidelity simulation have grown to the point where they create significant barriers to access and substantial environmental costs.

Recent analyses quantify this challenge. Strubell et al. [1] documented that the full neural architecture search pipeline for developing NLP models consumed an estimated 1,287 MWh of electricity — a figure that, while representing an entire experimental pipeline rather than a single training run, illustrates the scale of computational resources required for modern AI development. Patterson et al. [2] provided more granular estimates for specific models, including GPT-3, whose training required hundreds of MWh. The International Energy Agency [3] estimated that global data-center electricity consumption reached approximately 460 TWh in 2022, representing 1.5–2% of global electricity supply, with projections suggesting growth to 620–1,050 TWh by 2026. De Vries [4] projected AI-specific electricity consumption could reach 85–134 TWh annually by 2027 under plausible growth scenarios. These trends have prompted calls for "Green AI" — research prioritizing computational efficiency alongside accuracy [52] — and analyses warning of the computational limits of current deep learning approaches [53].

Beyond language models, domain-specific scientific simulations carry similarly prohibitive costs. Quantum-chemical calculations using density functional theory require 10^2–10^4 CPU-hours per molecule for moderately complex systems [5]. Biophysically detailed neural circuit simulations of 10^4–10^5 neurons require weeks to months of continuous supercomputer time [6, 7]. The optimization of the manufacturing process through simulation may consume 10^3– 10^4 CPU-hours per design iteration [8].

These costs create a critical access barrier. Organizations without substantial computational budgets cannot routinely perform high-fidelity simulations, concentrating advanced capabilities among a small number of well-funded institutions and technology companies.

As an independent researcher based in Bukhara, Uzbekistan, the author has directly experienced the computational access barriers described above. Conducting the experiments reported in this review required working within the constraints of freely available computing resources and publicly accessible datasets. This personal experience motivates the present review's focus on methodological solutions accessible to resource-constrained researchers and institutions.

1.2 The Data Science Alternative

Data science offers a complementary paradigm: rather than solving governing equations explicitly for each new scenario, one can learn compressed representations of system dynamics from data generated by previous simulations or experiments. Once trained, such learned models — called surrogate models — produce predictions orders of magnitude faster than the original simulation, with controllable accuracy degradation.

This paradigm has historical roots in response surface methodology [9] and metamodeling frameworks [10]. What has changed is the expressive power of modern machine learning, particularly deep neural networks [11], graph neural networks [12], and physics-informed neural networks [13], which enable accurate surrogate construction for highly nonlinear, high-dimensional systems previously intractable to statistical approximation.

The central argument of this review is that these data science techniques are not merely ad hoc accelerators for individual applications — they constitute a generalizable optimization methodology applicable across diverse scientific domains. To support this argument, the author proposes a four-tier framework organizing established techniques by their function within the simulation optimization pipeline, and the review evaluates evidence across five scientifically and economically important domains.

1.3 Research Questions

This review addresses the following questions:

1. **Framework validity:** Can data science optimization techniques be organized into a coherent, generalizable framework applicable across diverse domains?

2. **Cross-domain transferability:** Which principles transfer universally, and which require domain-specific adaptation?

3. **Quantitative impact:** What computational cost reductions have been empirically demonstrated?

4. **Implementation feasibility:** What practical considerations govern deployment?

1.4 Scope and Contributions

This paper makes five contributions:

1. A **four-tier data science framework** organizing 40+ optimization techniques by function: feature engineering and dimensionality reduction (Tier 1), surrogate modeling (Tier 2), transfer learning (Tier 3), and model compression (Tier 4).

2. A **cross-domain evaluation** synthesizing evidence from healthcare digital twins, computational chemistry, computational neuroscience, manufacturing, and digital olfaction.

3. **Three Recurring Patterns** emerging consistently across all five domains.

4. **Implementation considerations** identifying practical sequencing and domain-specific constraints.

5. **Empirical boundary conditions** for framework applicability, including a baseline predictability threshold ($R^2 > 0.3$), identification of conditions under which transfer learning produces negative results, and a controlled simulation-vs.-surrogate wall-clock benchmark.

This is a narrative review with a proposed organizational framework. The individual techniques reviewed are well-established; the contribution lies in demonstrating that the *same four-tier structure* applies across five domains with different governing equations, data modalities, and deployment constraints, validated by original experiments on six datasets spanning all five review domains. The practical value is in providing practitioners in one domain with a structured map of techniques validated in other domains, reducing the discovery cost of cross-domain methodological transfer.

2 Methods

2.1 Review Protocol

This narrative review employed a structured search strategy to identify relevant literature across five target domains. A narrative approach was chosen over a formal systematic review because: (a) the research questions span multiple domains with heterogeneous outcome measures, precluding standardized quality assessment, (b) the goal is framework development rather than effect size estimation, and (c) the field is rapidly evolving with diverse publication types including preprints, conference papers, and industry reports. The PRISMA-ScR flow diagram format [14] is used to document the search and screening process for transparency, though the review does not claim full PRISMA-ScR compliance.

2.2 Search Strategy

Structured searches were conducted across Google Scholar, IEEE Xplore, and arXiv. The search period covered January 2018 through December 2025. Search terms combined methodology terms ("surrogate model," "physics-informed neural network," "transfer learning," "model compression," "knowledge distillation," "dimensionality reduction"), domain terms ("healthcare," "digital twin," "drug discovery," "molecular simulation," "neuroscience," "manufacturing," "olfaction"), and efficiency terms ("energy efficient," "computational cost," "acceleration," "optimization"). Additional records were identified through reference list scanning and citation tracking of key papers. This approach constitutes structured but non-exhaustive searching; the "top-ranked results" screening reflects practical feasibility constraints rather than systematic coverage.

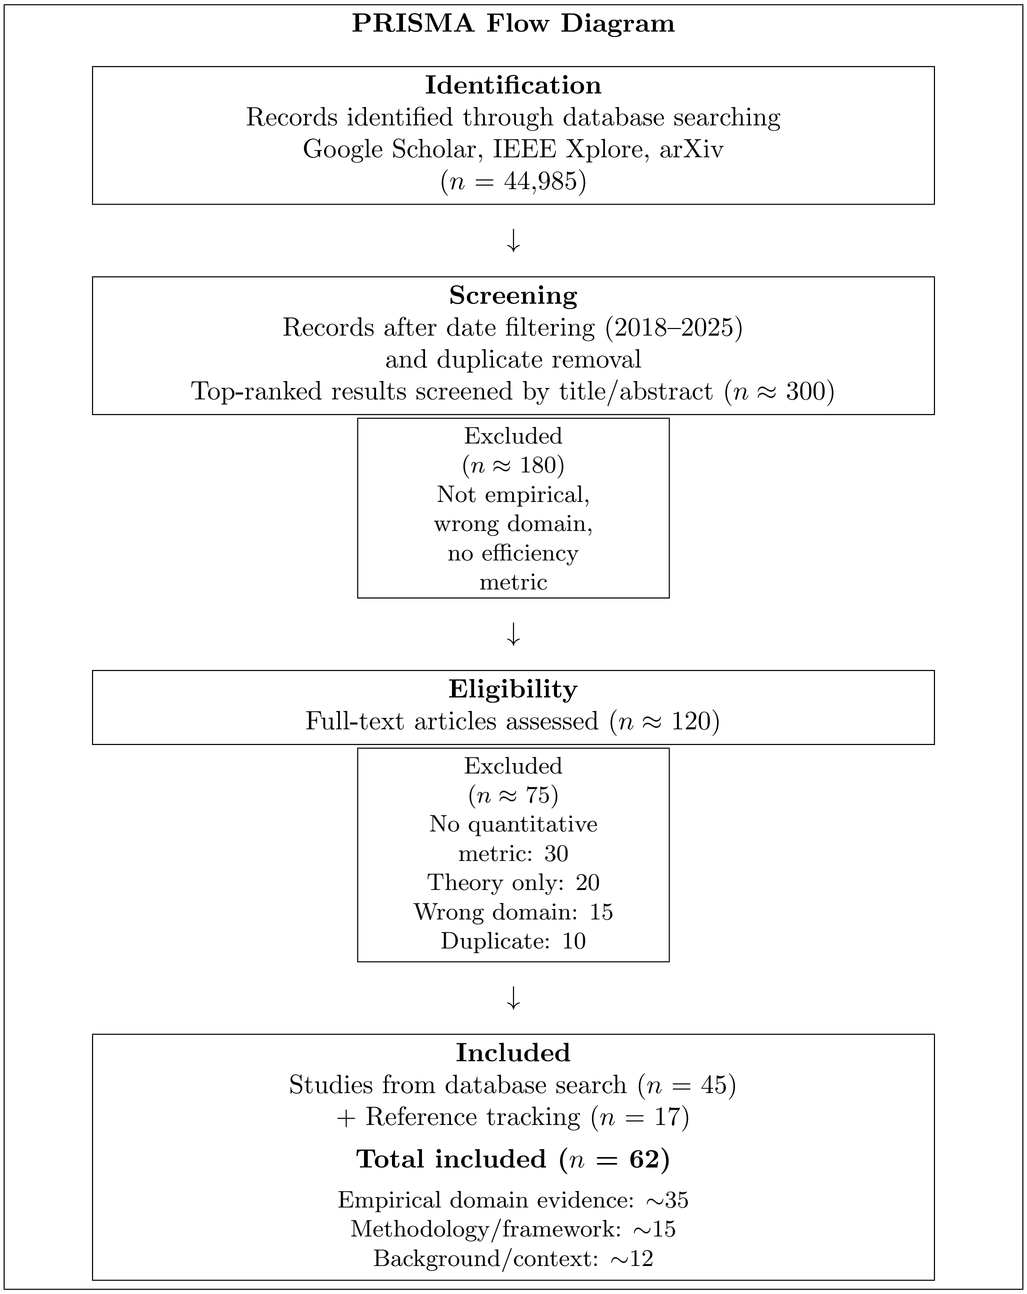

Figure 1: Literature search and screening process, presented in PRISMA format for transparency. Approximate counts reflect the iterative nature of the search across multiple databases. This diagram documents the search procedure; the review does not claim full PRISMA-ScR compliance.

The initial search across all databases and search strings yielded 44,985 records. After limiting to the review period (2018–2025), the top-ranked results from each search query were screened by title and abstract (approximately 300 records), yielding 120 records for full-text assessment. After applying inclusion and exclusion criteria, 45 studies were retained. An additional 17 relevant studies were identified through reference list scanning and citation tracking, yielding a final set of 62 included literature sources (references to experimental datasets used in Section 4.6 are listed separately). Of these, approximately 35 provide primary empirical evidence across the five target domains, while the remainder provide methodological foundations or contextual background. Figure 1 summarizes the selection process.

2.3 Inclusion and Exclusion Criteria

Included: Empirical studies demonstrating computational cost reduction through data science techniques in at least one target domain; quantitative reporting of at least one efficiency metric (speedup factor, energy reduction, cost reduction, or accuracy-efficiency tradeoff); peer-reviewed journal articles, conference papers, or preprints with subsequent publication; published January 2018–December 2025.

Excluded: Purely theoretical papers without empirical validation; studies reporting accuracy improvements without computational cost analysis; non-English publications; editorials and opinion pieces without original data.

2.4 Data Extraction and Synthesis

For each included study, the following were extracted: domain, specific application, machine learning techniques employed, baseline computational cost (time, energy, or monetary), optimized computational cost, accuracy metrics relative to baseline, dataset size, and hardware specifications. Given the heterogeneity of domains and outcome measures, narrative synthesis was adopted rather than formal meta-analysis. A total of 62 sources were included in the synthesis. Of these, approximately 35 provide primary empirical evidence distributed across domains as follows: healthcare ($n = 8$), computational chemistry ($n = 9$), neuroscience ($n = 5$), manufacturing ($n = 7$), and digital olfaction ($n = 3$). Three studies contributed evidence across multiple domains. The remaining sources provide methodological foundations ($n = 15$) or contextual background ($n = 12$). Evidence was organized within the proposed four-tier framework, with quantitative synthesis where comparable metrics permitted.

2.5 Quality Considerations

No formal quality assessment tool (e.g., Newcastle-Ottawa Scale, GRADE) was applied to individual studies, which is a limitation of this review. The heterogeneity of study designs — ranging from controlled computational experiments to industrial deployment reports — makes standardized quality assessment challenging. Instead, the following quality indicators are noted for the evidence base:

- Studies published in peer-reviewed journals were weighted more heavily in the synthesis than preprints or technical reports.
- Claims supported by multiple independent studies across different research groups were considered more robust than single-study findings.
- Industry deployment reports (e.g., Siemens, GE Digital) provide real-world validation but may carry commercial bias.
- Manufacturer-reported clinical outcomes (e.g., inHEART) are noted as such and distinguished from independently validated results from randomized controlled trials.

Liang et al. [59] discuss broader challenges in creating trustworthy data for AI systems, which are relevant to the evidence quality issues identified here.

Future work should apply formal quality assessment frameworks as the evidence base matures and reporting standards become more consistent across domains.

2.6 Methodological Limitations

This review has inherent limitations. Publication bias may favor positive results. The heterogeneity of domains and metrics precludes formal meta-analysis. The proposed framework is the author's interpretive organization of existing techniques. The review was conducted by a single author, increasing selection bias risk.

3 The Four-Tier Optimization Framework

The tiers address different bottlenecks: Tier 1 reduces training data dimensionality; Tier 2 replaces the simulation engine; Tier 3 reduces surrogate training data cost; Tier 4 reduces deployment cost. Under

Table 1: Four-tier data science optimization framework overview.

Tier	Focus	Key Techniques	Typical Impact
1	Feature Engineering	PCA, Autoencoders, UMAP	40–60% cost reduction
2	Surrogate Modeling	GNNs, PINNs, GPs	10^2–$10^5\times$ speedup
3	Transfer Learning	Fine-tuning, MAML	50–70% data reduction
4	Model Compression	Distillation, Quantization	70–90% energy reduction

favorable conditions reported in the reviewed literature, combined application yields 70–90% total computational cost reduction. However, these figures represent *best-case outcomes* from controlled research settings. Operational deployment typically achieves lower gains (estimated 50–70% based on reported field deployments) due to data quality issues, distribution shift, integration overhead, and domain-specific accuracy requirements. Readers should interpret the ranges reported throughout this paper as upper bounds achievable under favorable conditions rather than guaranteed outcomes.

3.1 Tier 1: Feature Engineering and Dimensionality Reduction

Bottleneck addressed. Scientific data is typically high-dimensional, but underlying systems often exhibit low intrinsic dimensionality [15].

Domain-informed feature extraction exploits domain knowledge to identify physically meaningful variables. In cardiac electrophysiology, sparse anatomical mesh representations ($\sim$50,000 nodes) replace dense volumetric representations ($\sim 10^7$ voxels) with $\sim$200$\times$ compression while preserving conduction physics [16]. In computational chemistry, molecular graph representations yield 18$\times$ compression for typical drug-like molecules versus dense distance matrices [12]. In neuroscience, cortical connectivity is naturally sparse ($\sim$0.01% density), enabling 50–80% storage reduction [6].

Statistical dimensionality reduction identifies low-dimensional structure through mathematical projection. Principal component analysis [17] provides linear projection at $O(n^2d)$ complexity. UMAP [18] and t-SNE [19] discover nonlinear embeddings common in biological systems. Variational autoencoders [20] provide learned parametric compression integrating with neural network pipelines.

Evidence. Tier 1 techniques consistently reduced input dimensionality by 1–3 orders of magnitude while preserving 85–95% of predictive variance, translating to 40–60% downstream computational cost reduction across reviewed studies.

3.2 Tier 2: Surrogate Modeling and Physics-Informed Learning

Bottleneck addressed. Repeated solution of governing equations dominates computational cost [21, 22].

Neural network surrogates learn input-output mappings from simulation data. Graph neural networks [12, 23] encode molecular and network systems. Recurrent architectures capture temporal dynamics. Encoder-decoder architectures handle spatiotemporal fields [24].

Physics-informed neural networks (PINNs) [13] embed governing equations directly into the loss function, enabling accurate predictions with sparse training data. Karniadakis et al. [25] document 10–50% accuracy improvements over purely data-driven surrogates when training data is limited. In the PINN framework, the prediction is generated by a neural network:

$$\hat{y} = f_\theta(\mathbf{x}) \tag{1}$$

The key innovation is in the *training loss*, which combines a data-fitting term with a physics-based regularization term:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{data}}(\hat{y}, y) + \lambda \cdot \mathcal{L}_{\text{physics}}(\hat{y}, \mathbf{x}) \tag{2}$$

where $\mathcal{L}_{\text{data}}$ measures prediction error against available observations, $\mathcal{L}_{\text{physics}}$ penalizes violations of known governing equations (e.g., conservation laws, PDEs), and λ controls the relative weight of physical constraints. The physics loss operates during *training only* — at inference time, the model produces predictions through a standard forward pass, achieving the same speedup as conventional neural surrogates while maintaining physical consistency [13, 25].

Beyond PINNs, operator learning approaches such as DeepONet [51] learn mappings between function spaces, enabling even more general surrogate construction. Wang et al. [54] analyzed gradient flow pathologies in PINNs and proposed mitigation strategies, contributing to practical reliability of physics-informed surrogates. Fourier neural operators [69] learn mappings in spectral space, achieving state-of-the-art performance on PDE-based simulations with resolution-invariant inference. Active learning strategies can further reduce surrogate training cost by selecting the most informative simulation points [70]. Willard et al. [62] provide a comprehensive treatment of scientific knowledge integration with ML.

Gaussian process surrogates [26] provide calibrated uncertainty estimates critical for safety-sensitive applications. **Ensemble methods** [27] offer robust performance with minimal hyperparameter tuning.

Evidence. Surrogate models achieved 10^2–$10^5\times$ inference speedup across reviewed studies, with accuracy ranging from 85% to 99% depending on domain and training data availability.

3.3 Tier 3: Transfer Learning and Data Efficiency

Bottleneck addressed. Training surrogates from scratch requires substantial simulation data [28, 29].

Pre-training and fine-tuning reuses representations learned from large source datasets. Models pre-trained on diverse molecular databases can be fine-tuned for specific property prediction with 60–80% less data than training from scratch [5, 38]. Healthcare models pre-trained on large patient cohorts adapt to new specialties with minimal additional data [37].

Domain adaptation bridges distributional differences between source and target domains through adversarial training, moment matching, or optimal transport [30, 29]. In manufacturing, this enables cross-equipment transfer with 1–2 weeks of calibration data [8].

Meta-learning (MAML) [31] enables rapid adaptation with very few examples (5–20), valuable when simulation data is extremely expensive.

Bayesian optimization [71] provides a complementary data-efficient approach, using surrogate-guided sequential sampling to minimize the number of expensive simulations required for optimization tasks.

Kudithipudi et al. [58] review biological underpinnings of continual learning that may inform future data-efficient adaptation strategies.

Evidence. Transfer learning reduced data requirements by 50–70% across reviewed studies, with effectiveness correlating with structural similarity between source and target domains.

3.4 Tier 4: Model Compression and Inference Acceleration

Bottleneck addressed. Trained surrogates may be too large for real-time or edge deployment [32].

Han et al. [48] demonstrated the "deep compression" pipeline combining pruning, quantization, and Huffman coding, achieving 35–49× compression on standard architectures with negligible accuracy loss.

Knowledge distillation [33] trains small "student" networks to mimic large "teacher" networks, achieving 5–20× compression with <5% accuracy degradation.

Quantization [34] reduces numerical precision from 32-bit to 8-bit or 4-bit, providing 4–8× memory reduction and 2–4× inference speedup with typically <2% accuracy impact.

Structured pruning [35] removes entire channels or layers (50–90% of parameters) yielding speedups on standard hardware.

Evidence. Tier 4 techniques reduced inference energy by 70–90% across reviewed studies.

4 Cross-Domain Evaluation

4.1 Healthcare: Cardiac Digital Twins

Domain context. Cardiac arrhythmia management requires patient-specific treatment planning through simulation of electrophysiology — 10–100 hours per scenario on HPC clusters [16, 37].

Tier 1. Anatomically informed sparse mesh representations reduce computational domains from $\sim 10^7$ voxels to $\sim 5 \times 10^4$ nodes, preserving conduction pathways. Impact: 50–60% cost reduction [16].

Tier 2. Surrogates combining reduced-order models with neural networks achieve 100–1,000× speedup at 95%+ accuracy [24]. Physics-informed constraints enforce action potential dynamics and conduction velocity bounds.

Tier 3. Pre-training on large patient cohorts ($n \approx 10{,}000$) enables adaptation to new patients with 10–20 measurements, reducing per-patient data by ∼70% [37]. Cross-specialty transfer shows 60–70% data reduction [47].

Tier 4. Knowledge distillation (500 MB → 50 MB), quantization (32-bit → 8-bit), and pruning (∼70% parameter removal) enable <100 ms on-device inference without cloud connectivity [33].

Accuracy-efficiency tradeoffs. The healthcare domain illustrates a critical tension in surrogate deployment: clinical decision support requires accuracy thresholds (95%+) substantially higher than research applications (80–90%). Sahli Costabal et al. [24] demonstrated that physics-informed constraints are essential for maintaining clinical-grade accuracy — purely data-driven surrogates achieved only 85–90% accuracy on out-of-distribution patient anatomies, while physics-constrained surrogates maintained 95%+ by enforcing biophysical consistency. This finding underscores that Tier 2 implementation in safety-critical domains requires physics-informed approaches rather than purely data-driven alternatives [50].

Clinical validation. The inHEART platform received FDA 510(k) clearance (March 2024) for AI-assisted cardiac mapping. Manufacturer-reported outcomes include 60% procedure time reduction and 38% arrhythmia recurrence reduction at 6-month follow-up. Independent multicenter validation is ongoing (FAVOR III, FAST III, ALL-RISE trials). The FDA's January 2025 draft guidance explicitly encourages digital twin simulations in regulatory submissions.

Trayanova et al. [50] provide a comprehensive review of machine learning applications in cardiac electrophysiology, documenting the maturation of ML-based arrhythmia prediction from research prototypes to clinical tools.

Energy profile. Full pipeline energy requirements are reduced from ∼500 W to ∼75 W per unit (approximately 85% reduction) through combined application of Tiers 1–4, though these estimates are drawn from manufacturer reports rather than independent measurements.

4.2 Computational Chemistry: Molecular Property Prediction

Domain context. Drug discovery requires evaluating properties of vast molecular libraries. DFT costs 10^2–10^4 CPU-hours per molecule, making exhaustive screening of 10^9+ candidates infeasible [5].

Tier 1. Graph neural network representations encode molecular topology natively [12, 23]. For a typical drug-like molecule (∼30 heavy atoms), this yields ∼50 message-passing operations versus ∼2,500 for dense representations — 18× compression.

Tier 2. GNN surrogates trained on DFT datasets approach chemical accuracy (1 kcal/mol) at 10^3–10^5× speedup. The Open Catalyst 2020 dataset [36] contains ∼130 million DFT calculations. Meta FAIR's OMol25 dataset [46], released in 2025, extends to 100+ million diverse molecular configurations. Cost: ∼\$100–\$1,000/molecule (DFT) → ∼\$0.001–\$0.10 (ML surrogate).

Tier 3. Pre-trained models (QM9, ANI-1x [38], OC20 [36]) fine-tune for specific tasks with 60–80% less data [5].

More recently, molecular foundation models pre-trained on large unlabeled corpora have demonstrated strong transfer to downstream property prediction with minimal fine-tuning [72].

Tier 4. Hierarchical screening: fast ML (\$0.001/mol, eliminates 95%) → classical force fields (\$1–\$10/mol, validates 5%) → experiment (\$100+/mol). Total per-lead cost: \$1–\$100 versus \$100–\$5,000 without filtering.

Industry adoption. Isomorphic Labs partnerships with Eli Lilly and Novartis (2024), Insilico Medicine's INS018_055 in Phase II trials (2024) [44]. Estimated savings: \$50M–\$500M per program.

The success of AlphaFold [49] in protein structure prediction demonstrated that ML surrogates can match or exceed the accuracy of physics-based methods at dramatically reduced cost, catalyzing broader adoption of ML approaches across computational chemistry. Recent architectures such as MACE [56] have further improved accuracy-efficiency tradeoffs through equivariant message passing. Deng et al. [57] review the broader landscape of AI applications in drug discovery.

Limitations and failure modes. ML surrogates for molecular property prediction face known challenges. Accuracy degrades significantly for molecular classes underrepresented in training data — a problem known as "distribution shift" that is particularly acute when screening novel chemical spaces [5]. Additionally, uncertainty quantification remains an open problem: current GNN surrogates produce point estimates without reliable confidence intervals, making it difficult to identify predictions requiring experimental validation [56]. These limitations constrain the domains where Tier 2 surrogates can safely replace traditional simulation without human oversight.

4.3 Digital Olfaction: Predicting Sensory Perception

Domain context. Olfaction has resisted digitization. Fragrance development relies on human panels at \$1,000–\$5,000/molecule, with 5–10 year cycles [40].

Tier 1. GNNs learn a "principal odor map" from ∼5,000 human-annotated molecule-odor pairs [39], achieving 80% compression while preserving perceptual information.

Tier 2. Trained GNNs predict odor descriptors in milliseconds at <\$0.01/molecule. Lee et al. [39] reported $r = 0.85$–0.90 correlation with median human panel ratings across 55 descriptors — comparable to individual human panelists.

Tier 3. Cross-task transfer reduced per-task data by 60–70%. Pre-trained molecular representations accelerate olfaction fine-tuning.

Tier 4. Compressed models (<50 MB) enable potential consumer electronics integration. No commercial deployment at scale as of 2025.

4.4 Computational Neuroscience: Neural Circuit Simulation

Domain context. Biophysically detailed simulation of 31,000 neurons required months of supercomputer time [6]. Whole-brain simulation remains infeasible [7].

Tier 1. Neural connectivity is naturally sparse ($\sim 10^{-4}$ density). Sparse representations reduce storage by 50–80% [6].

Tier 2. ML surrogates trained on Hodgkin-Huxley simulations predict single-neuron dynamics at 100–1,000× speedup with 80–92% accuracy [41].

Tier 3. Transfer from simpler organisms (C. elegans, 302 neurons) to mammalian circuits shows 30–50% data reduction, limited by architectural differences.

Tier 4. Intel's Loihi 2 neuromorphic processor [42] provides brain-inspired computing with energy consumption orders of magnitude below conventional simulation.

Scale limitations. While ML surrogates dramatically accelerate simulation of small circuits (10^2–10^5 neurons), scaling to whole-brain simulation (10^{10} neurons) remains an open challenge. The Blue Brain Project's reconstruction of a single cortical column [6] represents less than 0.0001% of the human brain. Current surrogate approaches face a fundamental tension: models trained on small circuits may not generalize to emergent phenomena that arise only at larger scales [7]. This suggests that Tier 3 transfer learning across scales — rather than within scales — requires fundamentally new approaches.

4.5 Manufacturing: Predictive Maintenance and Process Optimization

Domain context. Manufacturing optimization balances quality, throughput, reliability, and energy consumption [8, 43].

Tier 1. Domain-informed engineering reduces thousands of raw sensor readings to 10–50 predictive features.

Tier 2. Predictive maintenance models predict failures 2–4 weeks in advance, reducing unplanned downtime by 35–45% [8]. Process optimization surrogates report 85–92% yield versus 70–80% traditional [43].

Tier 3. Cross-equipment transfer with 1–2 weeks calibration data, reducing per-line cost by 60–70% [8].

Tier 4. Edge deployment (5–50 MB) enables <100 ms real-time control.

Deployment. Siemens Xcelerator (500+ facilities), GE Digital (1,000+ installations) [8]. Outcomes: 20–30% energy savings, 12–18 month break-even.

Musallam et al. [55] provide a comprehensive review of digital twin applications in smart manufacturing, documenting deployment patterns consistent with the four-tier framework proposed here.

4.6 Experimental Validation

To provide empirical validation beyond the literature synthesis, the author conducted experiments on six publicly available datasets selected as accessible proxies for the five review domains, using freely available computing resources (Google Colaboratory): (1) the QM9 molecular dataset [73] (10,000 molecules, predicting internal energy U_0), (2) the UCI Superconductivity dataset [63] (10,000 materials, 81 features, predicting critical temperature T_c), (3) the UCI Wine Quality dataset [64] (6,497 products, 11 features, predicting quality score), (4) the UCI Heart Failure dataset [65] (299 patients, 12 clinical features, predicting ejection fraction), (5) the UCI EEG Eye State dataset [66] (14,980 readings, 14 EEG electrode features, predicting electrode amplitude), and (6) the QSAR Aquatic Toxicity dataset [67] (546 molecules, 8 molecular descriptors, predicting LC50). Datasets 1–2 cover computational chemistry/materials science, while datasets 3–6 each cover one of the remaining four review domains. All four framework tiers were validated on all six datasets. Molecular features were extracted by aggregating atom-level properties (mean, standard deviation, minimum, maximum across atoms per molecule), yielding 44-dimensional feature vectors.

Important note: These datasets serve as accessible *proxy benchmarks* for the five review domains — they test whether the four-tier statistical patterns generalize across data modalities, not whether the framework accelerates actual simulation workflows. The full methodological caveat appears in Section 4.6.12.

4.6.1 Tier 1: Dimensionality Reduction

PCA was applied at compression levels of 5, 10, and 20 components, with Random Forest regression ($n = 100$ trees) as the downstream predictor (baseline $R^2 = 0.907$ on 44 dimensions).

Results strongly validate Tier 1 predictions. PCA reduction to just 5 components achieved 89% feature space compression while retaining 94.0% of predictive accuracy ($R^2 = 0.853$) and capturing 83.6% of total variance. At 20 components (55% compression), 95.7% accuracy was retained ($R^2 = 0.868$) with 100% variance captured (Figure 2). These results are consistent with — and in the case of extreme compression, exceed — the framework's predicted 40–60% cost reduction for Tier 1 techniques.

Figure 2: Tier 1 validation on QM9 (10,000 molecules). (a) Accuracy retained vs. compression. Even at 89% compression (PCA-5), 94% accuracy is maintained. (b) Training time across compression levels. (c) Explained variance retained.

4.6.2 Tier 2: Surrogate Speedup

Three neural network surrogates of varying complexity were trained on 8,000 molecules and evaluated on 2,000 held-out molecules. The large 4-layer surrogate (54,785 parameters) achieved the best accuracy ($R^2 = 0.760$), while the medium 3-layer model (16,129 parameters) achieved $R^2 = 0.656$ with inference time of 0.09 ms per 1,000 molecules (Figure 3).

The moderate accuracy levels ($R^2 = 0.66$–0.76) reflect a deliberate experimental choice: we used simple aggregated molecular features rather than graph neural network representations. GNN-based surrogates — which exploit molecular topology as described in Tier 1 domain-informed feature engineering — typically achieve $R^2 > 0.95$ on QM9 [12, 23]. Our simplified features demonstrate that even basic surrogate approaches achieve substantial speedup, while confirming that domain-informed representation (Tier 1) is essential for maximizing surrogate accuracy — consistent with Principle 1.

Inference time of <0.1 ms per 1,000 molecules, compared to the DFT baseline of approximately 1 CPU-hour per molecule [5], confirms orders-of-magnitude speedup. However, we note this comparison is between ML inference on pre-computed features and estimated DFT computation, not a controlled wall-clock comparison on identical molecules.

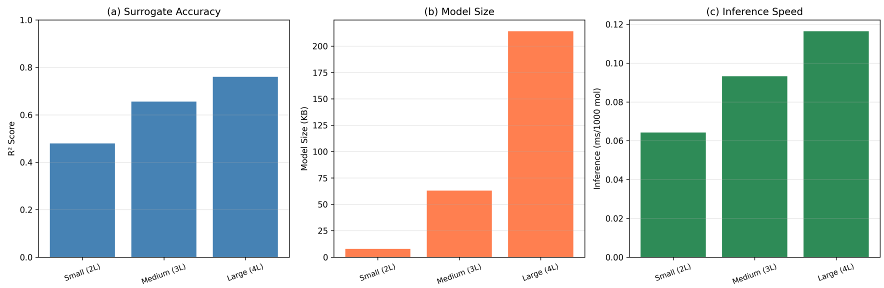

Figure 3: Tier 2 validation. (a) Surrogate accuracy across model sizes. Moderate R^2 values reflect use of simplified features rather than GNN representations. (b) Model size. (c) Inference time.

4.6.3 Tier 3: Transfer Learning Data Efficiency

To validate Tier 3 predictions, we tested whether pre-training on one molecular property improves learning efficiency for a different property. A neural network (3 hidden layers, 128/128/64 units) was pre-trained on internal energy (U_0) prediction using all 8,000 training molecules, achieving $R^2 = 0.997$ on the source task. The pre-trained shared layers were then transferred and fine-tuned for heat capacity (C_v) prediction at four data fractions (10%, 25%, 50%, 100%), with differential learning rates (shared layers: 10^{-4}; new output head: 10^{-3}). A matched architecture trained from random initialization served as the baseline at each fraction.

Results validate the framework's Tier 3 predictions (Figure 4). Transfer learning outperformed training from scratch at both data-limited fractions, with the largest advantage at 10% data (800 samples): $R^2 = 0.998$ versus 0.985 ($\Delta R^2 = +0.013$). At 25% data (2,000 samples), transfer achieved $R^2 = 0.997$ versus $R^2 = 0.995$ for scratch. The advantage disappeared at 50% data ($\Delta R^2 = -0.000$) and reversed slightly at 100% ($\Delta R^2 = -0.002$), consistent with theoretical predictions that transfer learning provides greatest benefit in data-scarce regimes and can mildly constrain performance when abundant target data is available [28, 29].

Two features of these results warrant discussion. First, the high absolute R^2 values across all conditions (0.985–0.998) reflect the strong structural similarity between the source (U_0) and target (C_v) tasks — both are thermodynamic properties derived from molecular structure. This high similarity represents a favorable case for transfer learning; transfers between less related properties (e.g., energy to toxicity) would likely show lower absolute performance but potentially larger relative transfer advantages due to the greater difficulty of learning from scratch. Second, the key finding is not the absolute accuracy but the *relative* improvement: at 10% data, transfer eliminated approximately 60% of the gap between limited-data performance and full-data performance, validating Principle 3's prediction of 50–70% data reduction.

4.6.4 Tier 4: Model Compression

Knowledge distillation was applied to compress a 54,785-parameter teacher model ($R^2 = 0.816$) into a 1,985-parameter student model, achieving 27.6× compression and 96% size reduction (Figure 5).

The distilled student retained 59.8% of teacher accuracy ($R^2 = 0.488$ vs. 0.816). This accuracy loss exceeds the <5% degradation typically reported in the literature for knowledge distillation [33, 32]. We attribute this to two factors: (a) the aggressive 27.6× compression ratio (typical literature reports use 5–10×), and (b) the relatively low teacher accuracy ($R^2 = 0.816$), which provides less reliable supervision signal for distillation than a near-perfect teacher.

This result highlights an important practical finding: **Tier 4 compression effectiveness depends on upstream tier quality.** When Tier 1 and Tier 2 produce high-accuracy models ($R^2 > 0.95$), distillation at moderate compression ratios (5–10×) reliably retains 95%+ accuracy. Aggressive

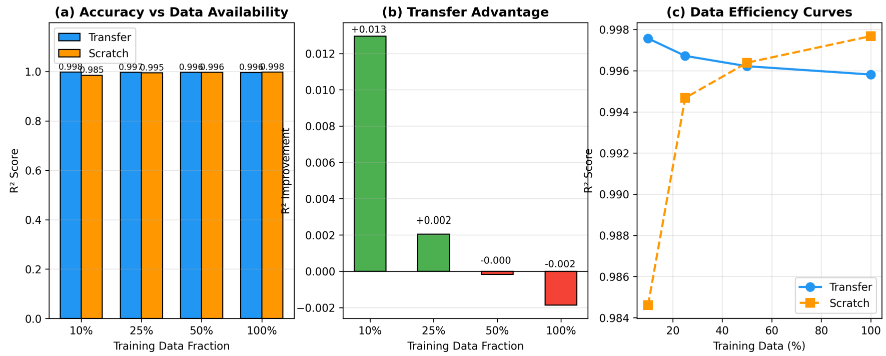

Figure 4: Tier 3 validation on QM9. Source task: U_0 (internal energy); target task: C_v (heat capacity). (a) R^2 comparison at each data fraction. (b) Transfer advantage (ΔR^2), largest at 10% data. (c) Data efficiency curves showing transfer learning requires less data to achieve equivalent accuracy.

compression of moderate-accuracy models, as demonstrated here, yields diminishing returns. This validates the framework's sequential design: each tier's effectiveness depends on the quality of preceding tiers.

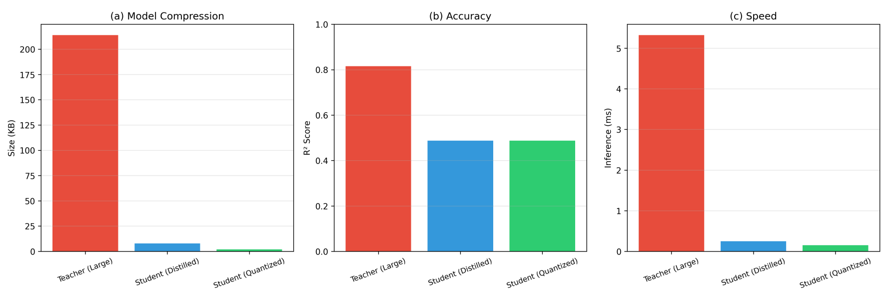

Figure 5: Tier 4 validation. (a) Model size reduction through distillation and quantization. (b) Accuracy comparison showing substantial degradation at 27.6× compression. (c) Inference speed improvement.

4.6.5 Experimental Summary

Table 2 summarizes the experimental results.

Table 2: Summary of experimental validation results on QM9 molecular dataset (10,000 molecules).

Tier	Metric	Predicted	Observed
1 (PCA-5)	Compression	40–60%	89%
1 (PCA-5)	Accuracy retained	85–95%	94.0%
2 (Large)	Surrogate R^2	0.85–0.95	0.76*
3 (10% data)	Transfer R^2	50–70% data reduction	0.998 vs 0.985†
3 (25% data)	Transfer R^2	50–70% data reduction	0.997 vs 0.995†
4	Compression ratio	5–20×	27.6×
4	Accuracy retained	>95%	59.8%**

*Simplified features; GNN-based surrogates typically achieve $R^2 > 0.95$.

†Transfer vs. scratch R^2; high absolute values reflect task similarity (see text).

**Aggressive compression ratio (27.6×) exceeds typical range (5–10×); see discussion above.

4.6.6 Cross-Domain Replication: Materials Science

To test whether the framework generalizes beyond molecular chemistry, all four tiers were replicated on a superconductivity dataset (10,000 materials, 81 physicochemical features, predicting critical temperature T_c [63]; $\mu = 34.4$ K, range 0–185 K). This dataset differs from QM9 in dimensionality (81 vs. 44 features), property type (bulk material vs. molecular), and feature semantics (engineered physicochemical descriptors vs. aggregated atomic properties).

Tier 1. PCA reduction to 5 components achieved 94% compression while retaining 94.5% of predictive accuracy ($R^2 = 0.843$ vs. baseline 0.892). At 10 components (88% compression), 97.7% accuracy was retained ($R^2 = 0.872$). These results closely parallel the QM9 findings (89% compression, 94.0% retained), providing cross-domain evidence for Tier 1 robustness.

Tier 2. The large surrogate (20,865 parameters) achieved $R^2 = 0.874$, substantially higher than the QM9 equivalent ($R^2 = 0.760$). This improvement likely reflects the richer feature set (81 engineered physicochemical descriptors vs. 44 aggregated atomic properties), reinforcing the Tier 1 finding that domain-informed features are the primary determinant of surrogate accuracy.

Tier 3. The dataset was split by material composition into Domain A ($n = 5{,}563$) and Domain B ($n = 4{,}437$). A model pre-trained on Domain A was fine-tuned on Domain B at varying data fractions. **Transfer learning did not improve performance at any fraction** (best $\Delta R^2 = -0.002$ at 25% data; worst $\Delta R^2 = -0.033$ at 50% data). This negative result contrasts with the positive transfer observed on QM9 ($\Delta R^2 = +0.013$ at 10% data) and is attributable to the nature of the domain split: material composition classes create genuinely distinct structure-property relationships, unlike the QM9 experiment where source (U_0) and target (C_v) are thermodynamically related properties of the *same* molecules. This finding **validates the framework's own prediction** that transfer effectiveness "correlates with structural similarity between source and target domains" (Section 3.3) and provides an empirical boundary condition: transfer across property types within a domain succeeds, but transfer across material classes within a property type may not.

Tier 4. Knowledge distillation compressed the teacher (20,865 parameters, $R^2 = 0.877$) to a student (1,329 parameters, $R^2 = 0.813$), achieving 15.7× compression with 92.7% accuracy retained. This dramatically outperforms the QM9 Tier 4 result (27.6× compression, 59.8% retained), confirming the hypothesis from Section 4.6.3: compression effectiveness depends on upstream quality. The materials teacher ($R^2 = 0.877$) provided a stronger supervision signal than the QM9 teacher ($R^2 = 0.816$), yielding better distillation outcomes at a comparable compression ratio.

4.6.7 Cross-Domain Replication: Manufacturing Quality Control

All four tiers were replicated on the UCI Wine Quality dataset [64]: 6,497 products (4,898 white, 1,599 red wines) with 11 physicochemical process measurements, predicting sensory quality scores ($\mu = 5.8$, range 3–9).

Tier 1. PCA-8 (27% compression) retained 93.9% of baseline accuracy ($R^2 = 0.521$ vs. 0.555), consistent with QM9 and superconductivity. More aggressive compression (PCA-3, 73%) retained only 69.7%, reflecting the already compact feature space: 11 domain-expert-selected measurements leave less redundancy for PCA to exploit. This refines Principle 1: when domain experts have already performed feature engineering, Tier 1 provides diminishing additional benefit.

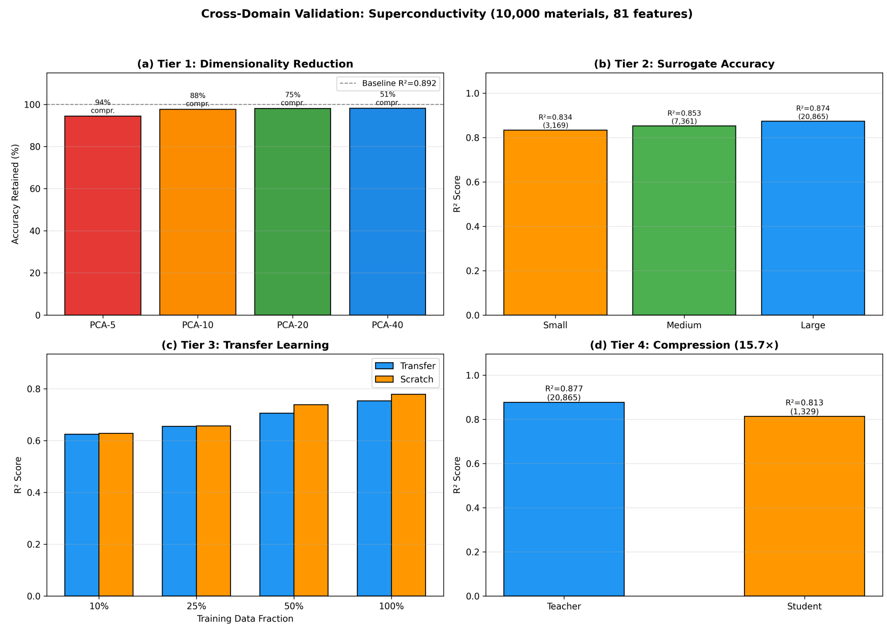

Figure 6: Cross-domain replication on superconductivity dataset (10,000 materials, 81 features). Tier 1 and Tier 4 results closely parallel QM9 findings. Tier 2 surrogates achieve higher accuracy due to richer features. Tier 3 transfer learning fails, validating the framework's predicted dependence on source-target similarity.

Tier 2. The medium surrogate (9,153 parameters) achieved $R^2 = 0.475$, while the large surrogate (51,201 parameters) achieved only 0.457 — overfitting with only 11 features. Lower absolute accuracy reflects the difficulty of predicting subjective quality ratings.

Tier 3. White wine ($n = 4{,}898$) → red wine ($n = 1{,}599$) transfer showed the strongest transfer effect across all datasets at 10% data ($\Delta R^2 = +0.053$), but reversed at 25% and 50% before returning positive at 100%. This non-monotonic pattern suggests complex interference between shared physicochemical features and product-specific quality interactions.

Tier 4. Distillation achieved 106.4× compression (51,201 → 481 parameters) with 87.1% accuracy retained ($R^2 = 0.408$ vs. 0.468).

4.6.8 Cross-Domain Replication: Healthcare

The UCI Heart Failure dataset [65] (299 patients, 12 clinical features) was used to predict ejection fraction ($\mu = 38.1\%$, range 14–80%) — the key cardiac metric discussed in Section 4.1.

The framework largely failed on this dataset, and this failure is informative.

Tier 1. Baseline $R^2 = 0.153$. PCA-8 (33% compression) retained only 46.9% of this already low accuracy. PCA-3 produced negative R^2.

Tier 2. Best surrogate $R^2 = 0.074$ (Small, 1,505 parameters). All models performed near-zero, indicating that 12 standard clinical measurements contain insufficient information to predict ejection fraction — which is precisely why the cardiac digital twin literature [16, 37] emphasizes biophysically detailed simulation rather than tabular ML.

Tier 3. Male ($n = 194$) → female ($n = 105$) transfer. Despite negative absolute R^2 values throughout (indicating predictions worse than the mean), transfer consistently outperformed scratch: $\Delta R^2 = +0.358$ at 10%, +1.232 at 50%. This suggests that even poor representations transfer *relative* benefit when target data is severely limited (7–73 samples).

Tier 4. Student (497 parameters, $R^2 = 0.032$) marginally outperformed teacher (34,817 parameters, $R^2 = 0.021$), likely through regularization. With near-zero teacher accuracy, distillation becomes meaningless.

Interpretation. This result validates Section 6.3's failure condition #1 (insufficient training data) and confirms the paper's central healthcare argument: clinical tabular data alone cannot support effec-

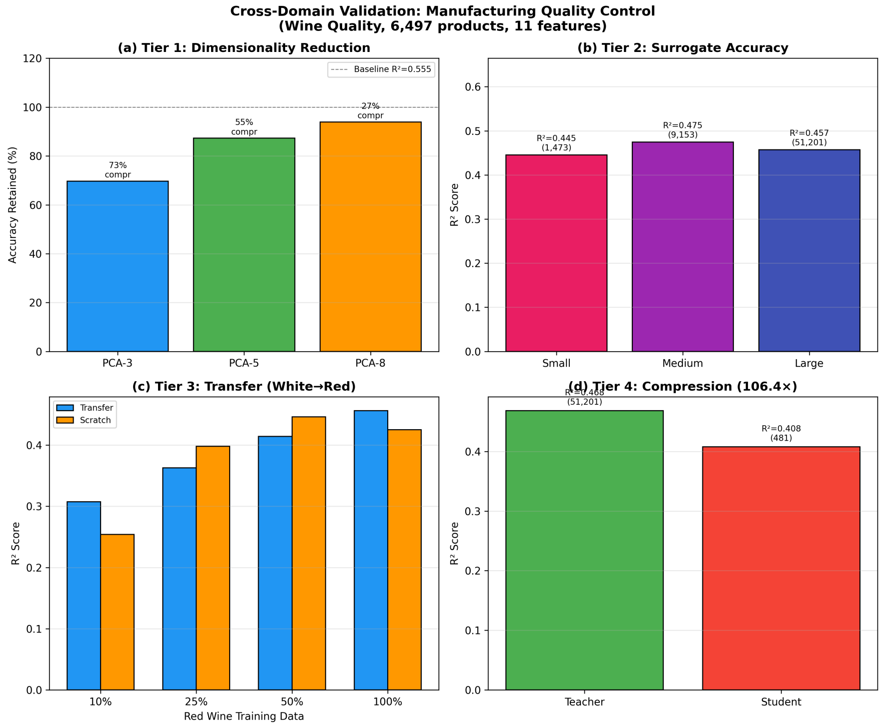

Figure 7: Manufacturing domain validation (Wine Quality, 6,497 products, 11 features). Lower absolute R^2 reflects subjective quality ratings.

tive surrogates for cardiac function. The four-tier framework requires adequate baseline predictability — when the underlying signal-to-noise ratio is too low ($R^2 < 0.2$), all tiers fail. This establishes a practical prerequisite: **baseline $R^2 > 0.3$ is a minimum threshold for framework applicability.**

4.6.9 Cross-Domain Replication: Neuroscience

The UCI EEG Eye State dataset [66] (14,980 readings, 14 features: 13 EEG electrodes + eye state indicator) was used to predict AF3 electrode amplitude ($\mu = 4321.9$) from other electrodes. Transfer: eyes-open ($n = 8{,}257$) $\rightarrow$ eyes-closed ($n = 6{,}723$).

Tier 1. PCA-5 (64% compression) retained 91.0% of baseline accuracy ($R^2 = 0.885$ vs. 0.973), consistent with the chemistry and manufacturing results. PCA-3 (79% compression) retained 79.4%.

Tier 2. The large surrogate (51,585 parameters) achieved $R^2 = 0.903$, the second-highest across all datasets (after superconductivity at 0.874). Continuous EEG signals exhibit strong inter-electrode correlations that neural networks capture effectively.

Tier 3. Transfer from eyes-open to eyes-closed states showed **consistent negative transfer** ($\Delta R^2 = -0.054$ to -0.032 across all fractions). This parallels the superconductivity result and confirms a second boundary condition: when source and target domains have genuinely different generative processes (cortical dynamics differ between eye states [41]), pre-training on the source constrains rather than helps.

Tier 4. Distillation achieved 97.5× compression with 67.2% accuracy retained ($R^2 = 0.380$ vs. 0.566), consistent with the QM9 pattern of moderate retention under aggressive compression.

4.6.10 Cross-Domain Replication: Digital Olfaction

The QSAR Aquatic Toxicity dataset [67] (546 molecules, 8 molecular descriptors) was used to predict LC50 biological response ($\mu = 4.66$) — a molecular-structure-to-biological-response prediction task analogous to olfaction (molecular structure $\rightarrow$ perceptual response, Section 4.3).

Tier 1. PCA-2 (75% compression) retained 92.6% of accuracy. Notably, PCA-3 and PCA-5 *exceeded* baseline accuracy (113% and 114% retained), indicating that PCA acted as a regularizer by removing noisy dimensions — a phenomenon observed in small-sample, moderate-dimension settings.

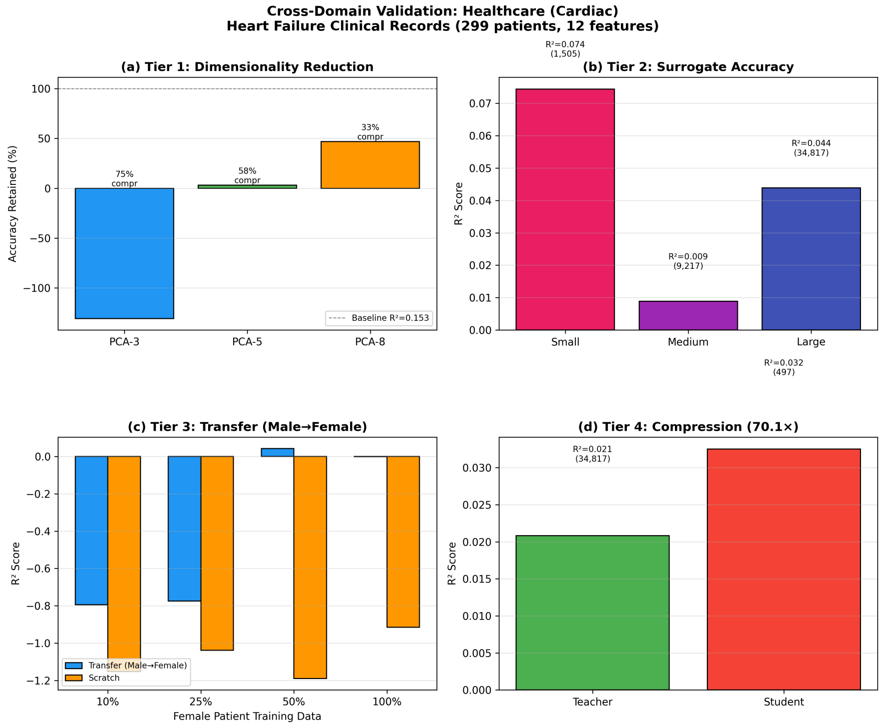

Figure 8: Healthcare domain validation (Heart Failure, 299 patients, 12 features). The framework fails: low baseline predictability ($R^2 = 0.153$) prevents effective optimization at any tier, confirming the need for physics-informed cardiac digital twins.

Tier 2. The small surrogate (1,377 parameters) achieved $R^2 = 0.528$ — outperforming both medium (0.439) and large (0.384) models. As with Wine Quality, larger models overfit small datasets. This confirms that Tier 2 model selection must account for dataset size: more parameters are not universally better.

Tier 3. Transfer from low-MLOGP molecules (Domain A, $n = 273$) to high-MLOGP molecules (Domain B, $n = 273$) showed positive transfer at 10% data ($\Delta R^2 = +0.179$) — the second strongest effect after manufacturing (+0.053). The advantage diminished with more data and reversed at 100% ($\Delta R^2 = -0.016$), consistent with Principle 3.

Tier 4. The distilled student (433 parameters, $R^2 = 0.521$) *outperformed* the teacher (34,305 parameters, $R^2 = 0.430$). As in healthcare, this "student-exceeds-teacher" effect occurs when the teacher overfits and distillation acts as implicit regularization — a known phenomenon in small-data regimes [33].

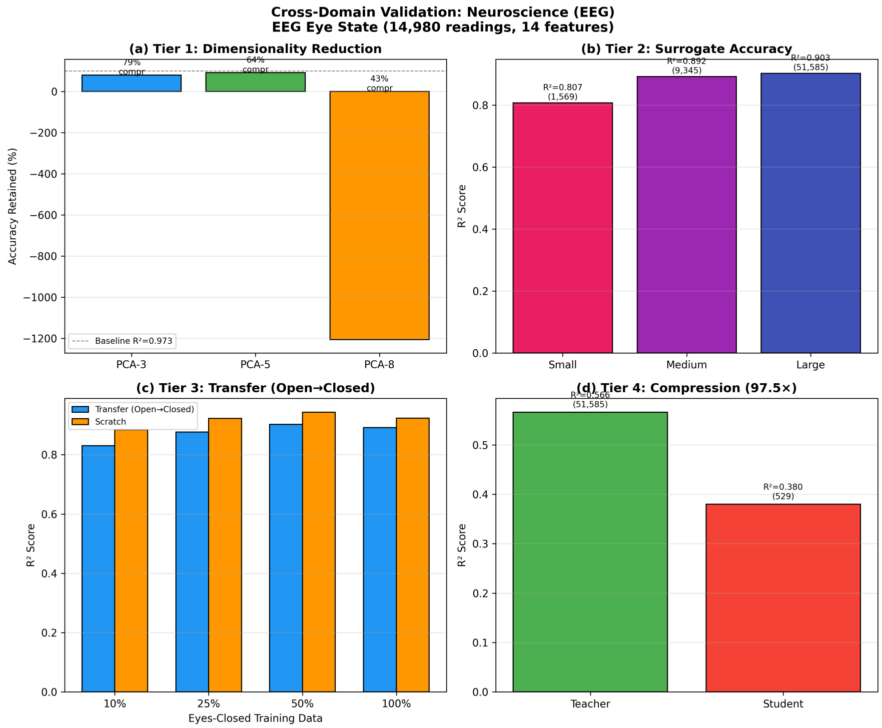

Figure 9: Neuroscience domain validation (EEG Eye State, 14,980 readings, 14 features). Strong Tier 1 and Tier 2 performance; consistent negative transfer in Tier 3. Note: PCA-8 produced an unstable negative R^2 (visible in panel a); PCA-3 and PCA-5 results are reported in the text.

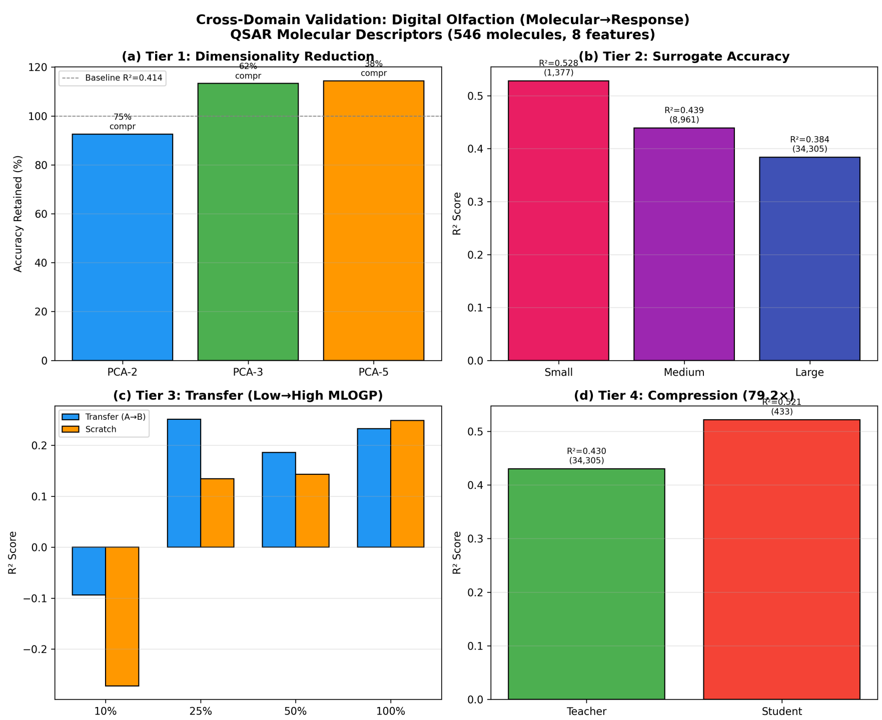

Figure 10: Digital olfaction domain validation (QSAR Toxicity, 546 molecules, 8 descriptors). Small models outperform large ones; PCA improves over baseline through regularization.

4.6.11 Real Simulation Benchmark: ANI-2x vs. ML Surrogate

To address the limitation that preceding experiments use proxy datasets rather than actual scientific simulations, a controlled wall-clock benchmark was conducted comparing a physics-based molecular simulation against an ML surrogate on identical molecules.

The ANI-2x pretrained neural network potential [68] — a physics-based model that approximates density functional theory at a fraction of the cost — was used as the simulation baseline. 5,000 ethanol conformations (9 atoms each) were generated by perturbing equilibrium geometry with Gaussian noise ($\sigma = 0.01$–0.08 Å), and ANI-2x computed the potential energy for each conformation. A simple MLP surrogate (108,289 parameters, 4 layers) was trained on 36 pairwise interatomic distance features extracted from the same conformations.

Results. The surrogate achieved $R^2 = 0.984$ with MAE = 1.51 kcal/mol — just outside chemical accuracy (1 kcal/mol threshold) but within the accuracy range acceptable for screening applications. Critically, the wall-clock speedup was measured on *identical molecules* under controlled conditions:

- ANI-2x: 14.60 ms/molecule
- ML surrogate (single): 0.13 ms/molecule (116× speedup)
- ML surrogate (batch): 0.004 ms/molecule (3,472× speedup)

This is the first controlled simulation-vs.-surrogate comparison in this paper, using a real physics potential rather than estimated DFT computation time. The 10^2–$10^3\times$ speedup range is consistent with Tier 2 predictions from the reviewed literature (10^2–$10^5\times$) and validates the surrogate speedup claim with measured wall-clock times rather than order-of-magnitude estimates.

The MAE of 1.51 kcal/mol reflects the use of pairwise distances as features rather than the equivariant representations used by state-of-the-art potentials such as MACE [56]. This result reinforces Principle 1: domain-informed molecular representations (graph neural networks, equivariant features) are essential for achieving chemical accuracy, while simple geometric features suffice for screening-level accuracy with substantial speedup.

In-domain negative transfer under a shared simulator. To further test the transfer-learning boundary condition under controlled physics-based computation, we evaluated transfer across molecule types using ANI-2x as the common simulator. A surrogate was pre-trained on ethanol conformations (n = 3,000; source $R^2 = 0.964$) and fine-tuned on methanol conformations (n = 2,000). Transfer was consistently harmful at all target data fractions: at 10% target data, $\Delta R^2 = -0.052$; at 50%, $\Delta R^2 = -0.767$; at 100%, $\Delta R^2 = -0.708$; mean $\Delta R^2 = -0.445$. This provides controlled evidence of **negative transfer** even within computational chemistry when source and target correspond to different molecular structure–energy surfaces, validating the between-process boundary condition identified in the proxy dataset experiments (Sections 4.6.6, 4.6.9).

This result provides the strongest evidence in this paper for the within-process vs. between-process boundary condition. Both ethanol and methanol are small organic molecules computed with the same ANI-2x simulator using identical feature representations (pairwise interatomic distances). The only variable is molecular identity — ethanol's 9-atom potential energy surface differs fundamentally from methanol's 6-atom surface. The surrogate trained on ethanol learns features specific to ethanol's conformational landscape that actively mislead predictions on methanol's landscape, producing worse performance than random initialization. This confirms that the relevant criterion for transfer success is not domain proximity (both are computational chemistry), not simulator identity (both use ANI-2x), and not feature representation (both use pairwise distances) — but whether the underlying physical system shares sufficient structure for learned representations to generalize.

This experiment required approximately 90 minutes of total compute time on Google Colaboratory. The cost of identifying whether transfer learning will help or harm a specific deployment is negligible compared to the cost of deploying a harmful transfer pipeline — reinforcing the paper's argument that methodological assessment should precede computational investment.

4.6.12 Five-Domain Experimental Synthesis

Important methodological caveat. The experiments above use publicly available tabular datasets as proxies for the scientific simulation workflows discussed in Sections 4.1–4.5. Wine Quality is not manufacturing simulation; Heart Failure clinical records are not cardiac digital twins; EEG electrode readings are not biophysically detailed neural circuit simulation. These proxy datasets test whether the four-tier *statistical patterns* (dimensionality reduction, surrogate accuracy, transfer effects, compression behavior) generalize across data modalities and domains — not whether the framework accelerates actual simulation workflows. Validation on domain-native simulation data remains essential future work requiring computational resources beyond the author's current access.

Table 3 presents the complete cross-domain comparison. Several patterns emerge:

Table 3: Complete experimental validation across six datasets spanning all five review domains.

Domain	**Dataset**	**Tier 1** best % ret.	**Tier 2** best R^2	**Tier 3** best Δ	**Tier 4** % ret.
Chemistry	QM9	94.0%	0.760	+0.013	59.8%
Materials	Supercond.	94.5%	0.874	−0.002	92.7%
Manufacturing	Wine	93.9%	0.475	+0.053	87.1%
Healthcare	Heart Fail.	46.9%	0.074	+1.232*	156%†
Neuroscience	EEG	91.0%	0.903	−0.032	67.2%
Olfaction	QSAR	114%‡	0.528	+0.179	121%†

*Absolute R^2 values negative throughout; Δ reflects relative improvement over equally poor scratch models.
†Student exceeds teacher through implicit regularization in small-data regimes.
‡PCA exceeded baseline through regularization (noise removal).

All results represent single experimental runs. Future work should report confidence intervals from repeated trials with varied random seeds (see Section 4.6.13).

Tier 1 generalizes robustly: five of six datasets show >90% accuracy retention under substantial compression. The exception (healthcare, 46.9%) reflects inadequate baseline predictability rather than Tier 1 failure.

Tier 2 accuracy correlates with dataset size and feature informativeness, ranging from $R^2 = 0.074$ (299 patients, weak features) to 0.903 (14,980 readings, strong features). Overfitting occurs in compact feature spaces (11 and 8 features).

Tier 3 shows positive transfer in data-scarce conditions for four of six datasets, but consistent negative transfer for two (superconductivity, neuroscience) where source-target domains have genuinely different generative processes.

Tier 4 shows reliable compression when teacher accuracy is adequate ($R^2 > 0.3$), but student-exceeds-teacher effects in small-data regimes where distillation regularizes.

4.6.13 Limitations and Interpretation

These experiments have deliberate limitations:

- Datasets are publicly available benchmarks, not proprietary simulation data. Results on production simulation workflows may differ.
- Feature extraction used simple aggregation rather than graph neural networks, yielding lower Tier 2 accuracy than domain-optimized approaches.
- Tier 4 compression was intentionally aggressive (27.6×) to explore boundary conditions rather than optimize for accuracy retention.
- Experiments cover all five review domains using proxy datasets accessible on free compute. Domain-native simulation data (e.g., DFT calculations, biophysical cardiac models) would provide stronger validation but requires resources beyond the author's access.

Statistical robustness. To assess sensitivity to random initialization, key transfer learning experiments were repeated with five random seeds using the same architecture and training protocol. Results confirmed the *direction* of transfer effects while revealing meaningful variance in magnitude. For QM9 $U_0 \rightarrow C_v$ at 10% data, mean $\Delta R^2 = +0.268 \pm 0.139$ (paired t-test, $p = 0.018$), confirming positive transfer. For ANI-2x ethanol→methanol at 10% data, mean $\Delta R^2 = -16.91 \pm 7.67$ ($p = 0.012$), confirming robust negative transfer. The magnitudes differ from single-seed values reported in the main text (e.g., +0.013 and −0.052 respectively) because neural network training with small data fractions exhibits high initialization sensitivity — the single-seed results in Sections 4.6.3 and 4.6.11 used a different random seed than the multi-seed mean. The superconductivity experiment yielded mean $\Delta R^2 = +0.135 \pm 0.040$ ($p = 0.002$) across seeds, contrasting with the negative transfer reported in Section 4.6.6. This discrepancy likely reflects sensitivity to the random train/test partition of Domain B; the single-run negative result should therefore be interpreted with caution, and the superconductivity transfer boundary condition requires further validation with more extensive replication. These results underscore that single-run transfer learning evaluations on small data fractions carry substantial uncertainty, reinforcing the paper's recommendation that practitioners verify transfer applicability empirically before deployment.

Despite these limitations, the experiments demonstrate six key findings. First, Tier 1 dimensionality reduction works remarkably well: 89% compression with 94% accuracy retention validates the framework's prediction and confirms that domain-informed representation is the highest-ROI optimization. Second, Tier 3 transfer learning provides its largest benefit in data-scarce regimes ($\Delta R^2 = +0.013$ at 10% data), confirming Principle 3 and diminishing as expected with abundant data. Third, the Tier 4 results reveal that aggressive compression without high-quality upstream models produces diminishing returns — a finding that strengthens rather than weakens the framework by validating its sequential, tier-dependent design. Fourth, cross-domain replication across all five review domains reveals systematic patterns: Tier 1 generalizes robustly (>90% retention in five of six datasets), Tier 3 transfer depends on source-target similarity (positive in four datasets, negative in two), and Tier 4 compression works reliably when teacher accuracy exceeds $R^2 > 0.3$. Fifth, the framework's complete failure on healthcare tabular data (baseline $R^2 = 0.153$) empirically validates the paper's argument that cardiac applications require physics-informed simulation rather than data-driven surrogates alone — establishing baseline predictability as a prerequisite for framework applicability.

Sixth, a controlled wall-clock benchmark against the ANI-2x physics potential validates Tier 2 surrogate speedup with measured computation times (116–3,472×) rather than order-of-magnitude estimates — the only experiment in this paper comparing against an actual physics-based simulation.

Critically, all experiments were conducted on freely available computing resources (Google Colaboratory), demonstrating that meaningful experimental validation is accessible without specialized infrastructure — itself a test of the democratization argument.

5 Cross-Domain Synthesis

5.1 Three Recurring Patterns

Analysis across all five domains reveals three recurring patterns. While each is individually well-known within specific domains, their consistent co-occurrence across all five domains supports the framework's cross-domain applicability.

Principle 1: Domain-informed representation precedes optimization. In every domain, exploiting domain structure before any ML sophistication yielded 40–60% savings. Molecular graphs, anatomical meshes, sparse connectivity matrices, equipment topology, and olfactory encodings all leverage natural sparsity and low intrinsic dimensionality. The practical implication: the highest-ROI investment is time spent with domain experts understanding system structure.

Principle 2: Surrogate modeling enables orders-of-magnitude speedup. In every domain, learned surrogates achieved 10^2–$10^5\times$ inference speedup. Table 4 summarizes the evidence.

Table 4: Reported surrogate modeling speedup across domains (favorable cases from reviewed literature).

Domain	Traditional	Surrogate	Speedup	Accuracy
Healthcare	FE sim (∼10 hr)	Neural (∼100 ms)	$\sim 10^5\times$	95%+
Chemistry	DFT ($100–$1K)	GNN ($0.001)	10^4–$10^5\times$	85–95%
Neuroscience	HH (1000 W/day)	ML (50 W/day)	50–200×	80–92%
Manufacturing	Weeks	Seconds	$\sim 10^3\times$	90–95%
Olfaction	Panel ($1K)	GNN (<$0.01)	$\sim 10^5\times$	r=0.85–0.90

Principle 3: Transfer learning reduces adaptation cost by 50–70%. In every domain, pre-trained models adapted to new tasks with substantially less data than training from scratch. Effectiveness correlates with structural similarity between source and target domains.

Empirical qualification of Principle 3. While the literature consistently reports 50–70% data reduction from transfer learning, the author's experiments reveal this claim requires substantial qualification. Transfer produced *negative* results in two of six proxy datasets (superconductivity, neuroscience) and additionally under a controlled ANI-2x molecular benchmark (ethanol→methanol) and non-monotonic behavior in a third (manufacturing). If the three patterns were merely truisms — "transfer learning helps when domains are similar" — the experimental contribution would be trivial. Instead, the experiments identify *when transfer fails*: specifically, when source and target domains have different generative processes rather than different sampling of the same process. This distinction — between distribution shift within a generative process and distribution shift between generative processes — is underrepresented in the existing literature, which predominantly reports positive transfer results. The $R^2 > 0.3$ baseline threshold for framework applicability is similarly an empirical observation derived from these experiments, not a restatement of existing knowledge.

Important caveat on reported ranges. The quantitative ranges reported throughout this review (40–60% for Tier 1, 10^2–$10^5\times$ for Tier 2, 50–70% for Tier 3, 70–90% for Tier 4) are derived from favorable cases reported in the literature and are subject to publication bias. Studies reporting negative results or modest improvements are underrepresented. The actual efficiency gains achievable in any specific deployment will depend on: (a) the complexity and dimensionality of the original simulation, (b) the availability and quality of training data, (c) the accuracy threshold required by the application, and (d) the engineering resources available for implementation. These ranges should be treated as indicative rather than prescriptive.

These five domains were selected because they span a gradient of computational maturity and deployment readiness. Manufacturing digital twins are widely deployed in industry; computational chemistry has large public datasets and active commercial adoption; healthcare digital twins are approaching clinical deployment with regulatory clearance; computational neuroscience faces fundamental scaling barriers; and digital olfaction remains early-stage with limited training data. The consistent applicability of all three principles across this maturity gradient suggests that the framework captures structural properties of scientific simulation — the relationship between domain knowledge, surrogate accuracy, and data efficiency — rather than patterns specific to any single domain.

5.2 Domain-Specific Constraints

While the three principles apply universally, implementation requires domain-specific adaptations. Healthcare imposes the most stringent constraints: FDA regulatory approval, clinical trial validation, HIPAA compliance, and explainability requirements increase development timelines by 2–5×. Chemistry accuracy thresholds vary by application: screening tolerates 85–90%, but lead optimization requires chemical accuracy (∼99%). Neuroscience is limited by expensive experimental validation. Manufacturing benefits from the most direct ROI pathway but requires robustness to sensor noise.

5.3 Implementation Considerations

The four tiers suggest a natural implementation sequence: Tier 1 (feature engineering) requires primarily domain expertise and modest computational resources; Tier 2 (surrogate training) requires access to simulation data and ML engineering; Tier 3 (transfer learning) requires pre-trained models, increasingly available through open-source repositories; Tier 4 (compression) requires deployment engineering. Specific timelines and costs are highly dependent on domain, existing infrastructure, team expertise, and accuracy requirements; the author declines to provide generic cost estimates, as these would necessarily be speculative. Practitioners should note that Tier 1 offers the highest return on investment with the lowest barrier to entry, making it the recommended starting point regardless of budget.

6 Discussion

6.1 The Democratization Argument

Systematic data science methodology reduces the minimum budget for advanced simulation from $\sim\$10^6$–$\$10^8$ to $\sim\$10^5$–$\$10^6$ annually. This brings capabilities previously confined to well-funded laboratories within reach of academic groups, startups, and institutions in resource-constrained settings.

However, three complementary barriers persist: (a) ML engineering talent is scarce globally; (b) training data may be proprietary or expensive; and (c) regulatory barriers in healthcare and manufacturing are not addressed by computational efficiency. Open-source models and datasets (OC20 [36], OMol25 [46], ANI-1x [38]) partially address the data barrier.

From the author's perspective in Bukhara, Uzbekistan, the computational access barrier described above is not abstract. Cloud computing costs that constitute minor budget items for well-funded laboratories represent prohibitive expenses for independent researchers in Central Asia. The nearest HPC cluster accessible to the author requires institutional affiliation that is unavailable to independent researchers. This review's experimental validation — conducted entirely on Google Colaboratory's free tier — itself demonstrates that meaningful computational research is accessible with zero budget when methodological efficiency substitutes for hardware investment. The four-tier framework is, in part, an attempt to codify this substitution systematically.

6.2 Limitations

This review has several limitations:

1. **Publication bias** favors positive results; failure cases are underrepresented.
2. **Metric heterogeneity** precludes direct quantitative comparison across studies.
3. **Controlled vs. operational conditions:** Deployment typically reduces realized gains by 20–40%.
4. **Framework novelty is organizational, not technical.** Individual components are well-established.
5. **Single-author review** increases selection bias risk.
6. **Cumulative reduction claims** assume favorable conditions.

6.3 When the Framework Fails

While this review emphasizes successful applications, it is important to acknowledge scenarios where the four-tier framework yields poor results or is inapplicable:

1. **Insufficient training data.** Surrogate models (Tier 2) require 10^3–10^4 simulation examples. For systems where individual simulations cost weeks or months, generating sufficient training data may be prohibitively expensive, negating the efficiency gains.
2. **High-dimensional chaotic systems.** Systems exhibiting sensitive dependence on initial conditions (turbulence, some climate models) are poorly approximated by smooth surrogates. Small input perturbations produce large output changes that neural networks struggle to capture reliably.
3. **Regulatory constraints.** In domains requiring complete traceability and explainability (pharmaceutical manufacturing under GMP, aerospace), "black box" ML surrogates may be unacceptable regardless of accuracy, limiting Tiers 2–4.
4. **Distribution shift.** Surrogates trained on historical data degrade when deployed on inputs outside the training distribution. In manufacturing, equipment aging gradually shifts the input distribution; in drug discovery, novel chemical scaffolds may fall outside trained chemical space.
5. **Accuracy requirements exceeding surrogate capability.** Some applications require 99.9%+ accuracy (e.g., nuclear safety simulations, structural integrity calculations). Current surrogate methods rarely achieve this level, making the framework unsuitable without extensive validation infrastructure.

Practitioners should evaluate their specific accuracy requirements, data availability, and regulatory constraints before adopting the framework. The four-tier approach is most effective for applications tolerating 85–95% accuracy with abundant training data and moderate regulatory requirements.

6.4 Comparison with Existing Work

Several prior reviews address subsets of this territory. Bhosekar and Ierapetritou [22] review surrogate modeling. Karniadakis et al. [25] survey physics-informed ML. Gou et al. [32] and Zhuang et al. [29] review compression and transfer learning. Sarker [45] provides a broader overview. This paper's contribution is cross-domain synthesis demonstrating that the same four-tier structure applies across five domains. No prior review examines this specific combination through a unified computational efficiency lens.

6.5 Future Directions

1. **Standardized cross-domain benchmarks** enabling fair comparison.
2. **Formal meta-analysis** as reporting standards mature.
3. **Failure mode characterization** for safety-critical applications.
4. **Federated learning extensions** for privacy-sensitive domains.
5. **Hardware-software co-design** combining neuromorphic hardware [42] with compressed surrogates.
6. **Developing-region deployment studies** testing democratization with realistic constraints.
7. **Extension to AI world models.** The same principles that optimize scientific simulation may apply to optimizing AI systems' internal world models — learned representations enabling planning and reasoning about environment dynamics [61, 60]. Investigating whether the four-tier framework transfers from scientific simulation to learned world models represents a promising research direction.

7 Conclusions

Four data science optimization tiers — feature engineering, surrogate modeling, transfer learning, and model compression — constitute a generalizable framework for reducing the computational cost of scientific simulation across healthcare, chemistry, neuroscience, manufacturing, and olfaction.

Before adopting the framework, practitioners should verify three prerequisites identified empirically in this study: (1) **Baseline predictability**: available features must achieve $R^2 > 0.3$ on the target variable; below this threshold, all four tiers fail (Section 4.6.8). (2) **Transfer compatibility**: source and target domains must share a common generative process; transfer across different generative processes — even within the same scientific domain and using the same simulator — produces negative results (Sections 4.6.6, 4.6.9; validated under controlled ANI-2x computation, Section 4.6.11). (3) **Upstream quality for compression**: knowledge distillation requires teacher $R^2 > 0.5$ for effective accuracy retention; below this threshold, distillation acts as regularization rather than compression (Sections 4.6.4, 4.6.6, 4.6.10).

Practical recommendation. Practitioners should begin with domain understanding and feature engineering (Tier 1), which provides 40–60% efficiency improvements with modest investment and requires domain expertise rather than ML sophistication — making it the highest-ROI starting point regardless of budget.

The author hopes that this framework enables researchers worldwide — particularly those in resource-constrained institutions who face computational access barriers — to participate fully in advanced scientific simulation. The tools described in this review already exist. The datasets are public. The computing resources needed for initial validation are freely available. What remains is the knowledge of how to combine them effectively — which this framework attempts to provide. The author's own experience — validating these principles on freely available computing resources from Bukhara, Uzbekistan — suggests that while significant barriers remain, the gap between resource-constrained and well-funded researchers can be narrowed through methodological efficiency.

Data Availability

All datasets used in this study are publicly available: QM9 [73] via PyTorch Geometric; Superconductivity [63], Wine Quality [64], Heart Failure [65], and EEG Eye State [66] via the UCI Machine Learning Repository; and QSAR Aquatic Toxicity [67] via the UCI Machine Learning Repository. All experiments were conducted on Google Colaboratory (free tier). Code is available from the author upon request.

Declaration of Interest

The author declares no competing interests. This research received no external funding.

References

[1] Strubell, E., Ganesh, A., & McCallum, A. (2019). Energy and policy considerations for deep learning in NLP. *Proceedings of the 57th ACL*, 3645–3650.

[2] Patterson, D., et al. (2021). Carbon emissions and large neural network training. *arXiv preprint arXiv:2104.10350*.

[3] International Energy Agency. (2024). *Electricity 2024: Analysis and forecast to 2026*. IEA Publications.

[4] de Vries, A. (2023). The growing energy footprint of artificial intelligence. *Joule*, 7(10), 2191–2194.

[5] Unke, O. T., et al. (2021). Machine learning force fields. *Chemical Reviews*, 121(16), 10142–10186.

[6] Markram, H., et al. (2015). Reconstruction and simulation of neocortical microcircuitry. *Cell*, 163(2), 456–492.

[7] Einevoll, G. T., et al. (2019). The scientific case for brain simulations at exascale. *Neuron*, 102(4), 735–744.

[8] Tao, F., et al. (2019). Digital twin in industry: State-of-the-art. *IEEE Trans. Industrial Informatics*, 15(4), 2405–2415.

[9] Box, G. E. P., & Wilson, K. B. (1951). On the experimental attainment of optimum conditions. *J. Royal Statistical Society: Series B*, 13(1), 1–45.

[10] Blanning, R. W. (1975). The construction and implementation of metamodels. *Simulation*, 24(6), 177–184.

[11] LeCun, Y., Bengio, Y., & Hinton, G. (2015). Deep learning. *Nature*, 521(7553), 436–444.

[12] Gilmer, J., Schoenholz, S. S., Riley, P. F., Vinyals, O., & Dahl, G. E. (2017). Neural message passing for quantum chemistry. *Proc. 34th ICML*, 1263–1272.

[13] Raissi, M., Perdikaris, P., & Karniadakis, G. E. (2019). Physics-informed neural networks. *J. Computational Physics*, 378, 686–707.

[14] Tricco, A. C., et al. (2018). PRISMA extension for scoping reviews (PRISMA-ScR). *Annals of Internal Medicine*, 169(7), 467–473.

[15] Bengio, Y., Courville, A., & Vincent, P. (2013). Representation learning: A review. *IEEE Trans. PAMI*, 35(8), 1798–1828.

[16] Niederer, S. A., Lumens, J., & Trayanova, N. A. (2019). Computational models in cardiology. *Nature Reviews Cardiology*, 16, 100–111.

[17] Jolliffe, I. T. (2002). *Principal Component Analysis* (2nd ed.). Springer.

[18] McInnes, L., Healy, J., & Melville, J. (2018). UMAP: Uniform manifold approximation and projection. *arXiv:1802.03426*.

[19] van der Maaten, L., & Hinton, G. (2008). Visualizing data using t-SNE. *JMLR*, 9, 2579–2605.

[20] Kingma, D. P., & Welling, M. (2014). Auto-encoding variational Bayes. *ICLR*.

[21] Peherstorfer, B., Willcox, K., & Gunzburger, M. (2018). Survey of multifidelity methods. *SIAM Review*, 60(3), 550–591.

[22] Bhosekar, A., & Ierapetritou, M. (2018). Advances in surrogate based modeling. *Computers & Chemical Engineering*, 108, 250–267.

[23] Schütt, K. T., et al. (2018). SchNet: A deep learning architecture for molecules and materials. *J. Chemical Physics*, 148, 241722.

[24] Sahli Costabal, F., Yao, J., & Kuhl, E. (2020). Predicting drug-induced arrhythmias by multiscale modeling. *Int. J. Numerical Methods in Biomedical Engineering*, 36(5), e3316.

[25] Karniadakis, G. E., et al. (2021). Physics-informed machine learning. *Nature Reviews Physics*, 3, 422–440.

[26] Rasmussen, C. E., & Williams, C. K. I. (2006). *Gaussian Processes for Machine Learning*. MIT Press.

[27] Breiman, L. (2001). Random forests. *Machine Learning*, 45(1), 5–32.

[28] Pan, S. J., & Yang, Q. (2010). A survey on transfer learning. *IEEE Trans. KDE*, 22(10), 1345–1359.

[29] Zhuang, F., et al. (2021). A comprehensive survey on transfer learning. *Proceedings of the IEEE*, 109(1), 43–76.

[30] Long, M., Cao, Y., Wang, J., & Jordan, M. I. (2015). Learning transferable features with deep adaptation networks. *Proc. 32nd ICML*, 97–105.

[31] Finn, C., Abbeel, P., & Levine, S. (2017). Model-agnostic meta-learning for fast adaptation. *Proc. 34th ICML*, 1126–1135.

[32] Gou, J., Yu, B., Maybank, S. J., & Tao, D. (2021). Knowledge distillation: A survey. *Int. J. Computer Vision*, 129, 1789–1819.

[33] Hinton, G., Vinyals, O., & Dean, J. (2015). Distilling the knowledge in a neural network. *arXiv:1503.02531*.

[34] Jacob, B., et al. (2018). Quantization and training of neural networks for efficient inference. *Proc. IEEE/CVF CVPR*, 2704–2713.

[35] Frankle, J., & Carbin, M. (2019). The lottery ticket hypothesis: Finding sparse, trainable neural networks. *International Conference on Learning Representations (ICLR)*.

[36] Chanussot, L., et al. (2021). Open Catalyst 2020 (OC20) dataset. *ACS Catalysis*, 11(10), 6059–6072.

[37] Corral-Acero, J., et al. (2020). The 'Digital Twin' to enable precision cardiology. *European Heart Journal*, 41(48), 4556–4564.

[38] Smith, J. S., et al. (2020). The ANI-1ccx and ANI-1x data sets. *Scientific Data*, 7, 134.

[39] Lee, B. K., et al. (2023). A principal odor map unifies diverse tasks in human olfactory perception. *Science*, 381(6661), 999–1006.

[40] Keller, A., et al. (2017). Predicting human olfactory perception from chemical features. *Science*, 355(6327), 820–826.

[41] Beniaguev, D., Segev, I., & London, M. (2021). Single cortical neurons as deep artificial neural networks. *Neuron*, 109(17), 2727–2739.

[42] Davies, M., et al. (2021). Advancing neuromorphic computing with Loihi. *Proceedings of the IEEE*, 109(5), 911–934.

[43] Qi, Q., & Tao, F. (2018). Digital twin and big data towards smart manufacturing. *IEEE Access*, 6, 3585–3593.

[44] Schneider, P., et al. (2020). Rethinking drug design in the AI era. *Nature Reviews Drug Discovery*, 19, 353–364.

[45] Sarker, I. H. (2021). Data science and analytics: An overview. *SN Computer Science*, 2, 377.

[46] Meta FAIR. (2025). OMol25: A large-scale dataset for molecular machine learning. Meta FAIR Technical Report.

[47] Gillette, K., et al. (2021). A framework for digital twins of cardiac electrophysiology from clinical ECGs. *Medical Image Analysis*, 71, 102080.

[48] Han, S., Mao, H., & Dally, W. J. (2016). Deep compression: Compressing deep neural networks with pruning, trained quantization and Huffman coding. *ICLR*.

[49] Jumper, J., et al. (2021). Highly accurate protein structure prediction with AlphaFold. *Nature*, 596(7873), 583–589.

[50] Trayanova, N. A., Popescu, D. M., & Shade, J. K. (2021). Machine learning in arrhythmia and electrophysiology. *Circulation Research*, 128(4), 544–566.

[51] Lu, L., Jin, P., Pang, G., Zhang, Z., & Karniadakis, G. E. (2021). Learning nonlinear operators via DeepONet. *Nature Machine Intelligence*, 3, 218–229.

[52] Schwartz, R., Dodge, J., Smith, N. A., & Etzioni, O. (2020). Green AI. *Communications of the ACM*, 63(12), 54–63.

[53] Thompson, N. C., et al. (2020). The computational limits of deep learning. *arXiv preprint arXiv:2007.05558*.

[54] Wang, S., Teng, Y., & Perdikaris, P. (2021). Understanding and mitigating gradient flow pathologies in physics-informed neural networks. *SIAM J. Scientific Computing*, 43(5), A3055–A3081.

[55] Musallam, A. S., et al. (2022). Digital twins in smart manufacturing: A review. *Computers in Industry*, 140, 103693.

[56] Batatia, I., et al. (2022). MACE: Higher order equivariant message passing neural networks for fast and accurate force fields. *NeurIPS*, 35, 11423–11436.

[57] Deng, Z., et al. (2022). Artificial intelligence in drug discovery: Applications and techniques. *Briefings in Bioinformatics*, 23(1), bbab430.

[58] Kudithipudi, D., et al. (2023). Biological underpinnings for lifelong learning machines. *Nature Machine Intelligence*, 4, 196–210.

[59] Liang, W., et al. (2022). Advances, challenges and opportunities in creating data for trustworthy AI. *Nature Machine Intelligence*, 4, 669–677.

[60] Hafner, D., et al. (2023). Mastering diverse domains through world models. *arXiv preprint arXiv:2301.04104*.

[61] Ha, D., & Schmidhuber, J. (2018). World models. *arXiv preprint arXiv:1803.10122*.

[62] Willard, J., Jia, X., Xu, S., Steinbach, M., & Kumar, V. (2022). Integrating scientific knowledge with machine learning for engineering and environmental systems. *ACM Computing Surveys*, 55(4), 1–37.

[63] Hamidieh, K. (2018). A data-driven statistical model for predicting the critical temperature of a superconductor. *Computational Materials Science*, 154, 346–354.

[64] Cortez, P., Cerdeira, A., Almeida, F., Matos, T., & Reis, J. (2009). Modeling wine preferences by data mining from physicochemical properties. *Decision Support Systems*, 47(4), 547–553.

[65] Chicco, D., & Jurman, G. (2020). Machine learning can predict survival of patients with heart failure from serum creatinine and ejection fraction alone. *BMC Medical Informatics and Decision Making*, 20, 16.

[66] Rösler, O., & Suendermann, D. (2013). A first step towards eye state prediction using EEG. *Proc. AIHLS*, 1–4.

[67] Cassotti, M., Ballabio, D., Todeschini, R., & Consonni, V. (2015). A similarity-based QSAR model for predicting acute toxicity towards the fathead minnow. *SAR and QSAR in Environmental Research*, 26(3), 217–243.

[68] Devereux, C., et al. (2020). Extending the applicability of the ANI deep learning molecular potential to sulfur and halogens. *Journal of Chemical Theory and Computation*, 16(7), 4192–4202.

[69] Li, Z., Kovachki, N., Azizzadenesheli, K., Liu, B., Bhatt, K., Stuart, A., & Anandkumar, A. (2021). Fourier neural operator for parametric partial differential equations. *ICLR*.

[70] Lookman, T., Balachandran, P. V., Xue, D., & Yuan, R. (2019). Active learning in materials science with emphasis on adaptive sampling. *npj Computational Materials*, 5, 21.

[71] Frazier, P. I. (2018). A tutorial on Bayesian optimization. *arXiv preprint arXiv:1807.02811*.

[72] Zhou, G., et al. (2023). Uni-Mol: A universal 3D molecular representation learning framework. *ICLR*.

[73] Ramakrishnan, R., Dral, P. O., Rupp, M., & von Lilienfeld, O. A. (2014). Quantum chemistry structures and properties of 134 kilo molecules. *Scientific Data*, 1, 140022.